

Simulation (An Introduction)

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- Simulation: Coverage Probability of Confidence Intervals
- Simulation: Properties of Hypothesis Tests
- END

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- Simulation: Coverage Probability of Confidence Intervals
- Simulation: Properties of Hypothesis Tests
- END

Definition

\bar{X} : X_1, \dots, X_n iid \sim distrib with μ ; $\text{var} = \sigma^2$.
 $\Rightarrow \bar{X} \sim N\left(\mu; \frac{\sigma^2}{n}\right)$ if n is large enough.

- Simulation (in statistics) is using artificially generate data in order to test out a hypothesis or statistical method.
- Advantages of simulation: it is cheap; it is much faster than traditional data collection; in good conditions the results of simulation can approximate real results.
- Disadvantage: the results from simulation can approximate the real-results only, not the real results.

Example 1: The t -distribution (Theory)

$t_{df} \dots$

Definition

If $Z \sim N(0, 1)$ and $U \sim \chi^2_k$ and Z and U are independent, then the distribution of $Z/\sqrt{U/k}$ is called the t **distribution** with k degrees of freedom.

- Let X_1, \dots, X_n be a sample of n independent $N(\mu, \sigma^2)$ rv.
- Let $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ (sample mean).
- Let $S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$ (sample variance).
- Then, it can be proved that:

$$\frac{Z}{\sqrt{U/k}} = t$$

ST2131

$$\frac{\bar{X} - \mu}{S/\sqrt{n}} \sim t_{n-1}$$

Example 1: The t -distribution (Simulation)

- A sample of size $n = 9$ is taken from $N(0, \sigma^2)$.

- Consider the statistic

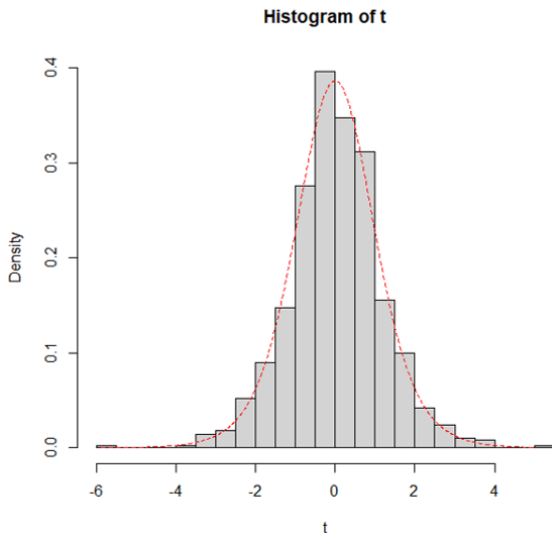
$$t = \frac{\bar{X} - \mu}{S/\sqrt{9}} \sim t_8$$

How do we “view” this result by simulation? The steps below can help.

- Generate N random samples (N can be any number, better to be a large one, like 1000), each sample of size $n = 9$ (fixed) from a normal distribution with mean 0 (fixed), and a (self random chosen) standard deviation σ .
- For each sample, compute the statistic t (and must save all the values/realizations of t).
- Construct a histogram for the N realizations of the statistic t .


The t -distribution (Simulation)

The histogram of N realizations of t and the plot of t_8 distribution (in red).



Example 2: Buffon's Needle Experiment

- A needle of length l is thrown randomly onto a grid of parallel lines with distance $d > l$.



- **Question:** What is the probability that the needle intersects a line?
- The answer is $\frac{2l}{\pi d}$. Can we get the answer through simulations?

Example 2: Buffon's Needle Experiment

Main idea of the solution by experiment:

- Simulate the throwing of a needle into a grid of parallel lines, say N times.
- Count the number of times the needle intersects a line, say n times.
- Then n/N gives an estimate of the probability that the needle intersects a line.
- A website gives the visualization of the experiment

<http://www.metablake.com/pi.swf>

Example 2: Buffon's Needle Experiment

How to get the solution by simulation?

- ➊ Generate the position and the inclination of the needle.
 - ▶ Generate a random number, x , from $U(0, d/2)$. This number represents the location of the center of the needle.
 - ▶ Generate a random number, q , from $U(0, \theta)$. This number represents the angle between the needle and the parallel lines.
- ➋ Check if the needle cuts a line.
 - ▶ The needle cuts a line if $x < l/2 \sin(\theta)$.
 - ▶ Create a variable $w = 1$ if $x < l/2 \sin(\theta)$ and 0 otherwise.
- ➌ Repeat Steps 1 and 2 a total of N times.
- ➍ Count the number of times $w = 1$, let say n times. An estimate of the probability of the needle intersecting a line is given by n/N .

Example 3: Comparing Estimators

sample mean ←
median ←
trimmed mean }
win. mean }

We have learnt about several robust estimators of location.

Question: Which estimator is better, the trimmed mean or the Winsorized mean?

Secondary Questions:

- Under what conditions (in terms of the underlying distributions) is the estimator better?
- What criteria to determine the performance of the estimator? Ans: in our course, we can use the Mean Square Error ($MSE = Bias^2 + Variance$).

It may be intractable to compute the (theoretical) MSE for each of the estimators under different underlying distributions.

Simulation is an alternative solution.

Question: How to design such a simulation study?

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- Simulation: Coverage Probability of Confidence Intervals
- Simulation: Properties of Hypothesis Tests
- END

A Set of Random Integer Numbers

- Computer can help us to choose a set of 10 random integer numbers, range from 1 to 100.

> sample(1:100,10,replace=T) #or

[1] 41 59 64 7 41 18 76 85 40 29

> sample(1:100,10,replace=F)

[1] 6 83 32 55 42 75 47 25 46 28

- That means, in the set of integers from 1 to 100, (1) it will randomly choose an integer. (2) if “replace = T”, it will put that number back to the series of 1 to 100 and then choose the second random number. This process continues till the 10th integer is chosen. This allows the output to have repeating integers.
- If “replace = F”, that means the chosen number will be removed from the series of 1 to 100, hence the output will have no repeating number.
- That sound simple!

Pseudo-random Numbers

Now, we want to generate a set of random numbers that follow Uniform(0,1) distribution. How to do that?

Definition

A sequence of uniform pseudo-random numbers $\{U_i\}$ is a deterministic sequence of number in $[0, 1]$ having the same relevant statistical properties as a sequence of random numbers.

- In R, we generate a set of 6 random numbers from $\sim \text{Unif}(0, 1)$ by

```
> runif(6, 0, 1)
```

rexp()

```
[1] 0.08988458 0.91487180 0.80281471 0.24061284 0.68319329 0.495
```

- In Python, we generate a set of 6 numbers that are iid $\sim \text{Unif}(0, 1)$ by

```
> # import numpy as np  
> # x = np.random.uniform(0, 1, 6)  
> # print(x)
```

- What's the logic behind when the random numbers above were generated?

Congruential Generators

- **Congruential generators** are defined by

$$X_{n+1} = (aX_n + c) \mod m$$

for a multiplier a , shift c , and modulus m . Here a, c and m are all integers.

- To initialize, we have to provide a seed ($n = 0$) which is X_0 .
If $c = 0$, generators having the form

$$X_{n+1} = aX_n \mod m$$

are called **multiplicative congruential generator**. If $c > 0$, they are called **linear congruential generator**.

- Uniform random numbers are obtained by

$$U_i = X_i/m$$

Some Properties of Congruential Generators

- The $m + 1$ values $\{X_0, X_1, \dots, X_m\}$ cannot be distinct and at least one value must occur twice, as X_i and X_{i+k} , say.

item $\{X_i, X_{i+1}, \dots, X_{i+k-1}\}$ is repeated as $\{X_{i+k}, X_{i+k+1}, \dots, X_{i+2k-1}\}$.

- The sequence $\{X_i\}$ is periodic with period $k < m$.
- For multiplicative generators, the maximal period is $m - 1$.
- If 0 ever occurs, it is repeated indefinitely.
- One of our primary objectives is to use a generator with as large period as possible, however that does not mean that a large period can help to guarantee a good generator.
- The values commonly used for parameters are: $m = 2^{32}$ or $m = 2^{31} - 1$, etc. $c = 1$ or $c = 0$ or $c = 12345$, etc. $a = 8121$ or $a = 22695477$, etc.

R function for Congruential Generators

- Creating a set of 10 random numbers that are from $Unif(0,1)$.

```
> cg <- function(n,a,m,c,x0){  
+ series <- NULL  
+ for (i in 1:n) {  
+ x1 <- (a*x0+c)%m  
+ x0 <- x1  
+ series <- c(series,x1/m) }  
+ return(series) }  
> cg(10,397204094,2^31-1,0,1234)  
[1] 0.24381116 0.08947511 0.38319371 0.72800325 0.72792771 0.175038  
[7] 0.40680994 0.90923700 0.34271140 0.55960111
```

Generate Uniform Random Numbers

We want to generate a set of n random numbers from $Unif(a, b)$ or $U(a, b)$ for short.

- In R:

runif(n,a,b).

- In Python:

import numpy as np

np.random.uniform(a,b,n) }

- In SAS:

data Ugen

call streaminit(1234); /* 1234 is used as the seed, this helps the generated random numbers reproducible */

do i = 1 to 10; $n = 10$

x = rand('uniform', 2, 3);

output;

end;

keep x;

run;

proc print data=Ugen;

var x;

run;

set. seed (999)

10 values ~ $U(2, 3)$

$a = 2$ $b = 3$

Generating Non-uniform Random Numbers: Inversion Method

cdf $F \rightarrow F^{-1}$

The theory of the inversion method is as below:

- If random variable X has a continuous distribution function (cdf) $F(x) = P(X \leq x)$ then a variable Y which is defined as $Y = F^{-1}(X)$ will follow $U(0, 1)$.
- If a variable $Y \sim U(0, 1)$ and variable X has cdf F then the cdf of the variable $F^{-1}(Y)$ is F .


That gives us a way to generate random numbers following a distribution function $F(x)$.

- Handwritten notes:* $X \sim \text{expo}$ has cdf: F → 10 values from exp (mean = 5) → $\text{rexp}(10, \text{rate} = 1/5)$
- 1 Generate U from $U(0, 1)$
 - 2 Set $X = F^{-1}(U)$ provided the inverse exists.
Then X is a random number from the distribution F .

Exponential Distribution

Theoretically,

- If X follows an exponential distribution with parameter λ (rate) (which then has mean $1/\lambda$), then its cdf is


$$F(x) = 1 - e^{-\lambda x}$$

- Let $U \sim U(0, 1)$, then $X = F^{-1}(U)$ will follow exponential distribution, where

$$F^{-1}(U) = -\lambda \log(1 - U)$$


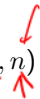

Hence, to generate a random number that follows exponential distribution with $\lambda = 5$, we'll:

- 1 Randomly generate U from $U(0, 1)$
- 2 Let $X = -5 \log(1 - U)$.

Then X is a random number from the exponential distribution with rate $\lambda = 5$.

Exponential Distribution

Generate n random numbers from $Exp(\lambda)$ where λ is the rate:

- In R, 
 $\text{rexp}(n, \lambda)$
- In Python,
import numpy as np
 $\text{np.random.exponential}(1/\lambda, n)$ 
- In SAS,
data Egen;
call streaminit(1234);
do i = 1 to 10; * here we use $n = 10$;
 $x = \text{rand}(\text{'exponential'}, 1/5);$ *rate $\lambda = 5$ 
 output;
end;
keep x;
run;
proc print data=Egen;
var x;
run;

Weibull Distribution (Standard version)

- If a variable X follows a Weibull distribution with one parameter of shape α with its pdf

$$f(x) = \alpha x^{\alpha-1} e^{-x^\alpha}, \quad x > 0,$$

then its cdf is


$$F(x) = 1 - e^{-x^\alpha}, \quad x > 0.$$

(The standard version has the scale parameter fixed at 1)

- The inverse function of $F(x)$ is $F^{-1}(x) = (-\log(1 - x))^{1/\alpha}$.
- If we want to generate a random number that follows Weibull distribution, we can:
 - 1 Randomly generate U from $U(0, 1)$
 - 2 Let $X = (-\log(1 - U))^{1/\alpha}$.
Then X is a random number from the standard Weibull distribution with shape α .

Generating a Weibull Random Variable

We'll generate 10 random numbers that follow Weibull distribution with shape $\alpha = 4$.

- In R,
`x = rweibull(10, shape = 4)`
- In Python,
`import numpy as np
x = np.random.weibull(4, 10)
print(x)`
- In SAS,
`data Weibullgen;
call streaminit(1234);
n = 10;
do i = 1 to n;
x = rand('weibull',4); 
output; end;keep x;run;`

Generating a Normal Random Variable

cdf of Normal
pdf: $\frac{1}{\sigma\sqrt{2\pi}} \exp\left\{-\frac{(x-\mu)^2}{2\sigma^2}\right\} \dots$

- Normal distribution does not have cdf with explicit form, hence we cannot apply the inversion method to generate a random number following normal distribution from a uniform number.
- A random number that follows standard normal distribution is generated by an indirect method as below.
 - Generate u_1 and u_2 from $U(0, 1)$.
 - Set $\theta = 2\pi u_1$.
 - Set $R = (-2 \log u_2)^{1/2}$.
 - Set $X = R \cos(\theta)$ and $Y = R \sin(\theta)$.
 - X and Y are independent standard normal variables.
 - For simplicity, only X or Y is used.

Generating a Normal Random Variable

To generate n random numbers that follow normal distribution with mean μ and standard deviation σ , we do:

- In R

```
n = 10; mu = 5; sigma = 1.5  
x = rnorm(n, mean = mu, sd = sigma)
```

- In Python,

```
import numpy as np  
n = 10  
mu = 5  
sigma = 1.5  
x = np.random.normal(loc = mu, scale = sigma, size = n)
```

- In SAS,

```
data Normalgen;  
call streaminit(1234);  
n = 10; mu = 5; sigma = 1.5;  
do i = 1 to n  
→ x = rand('normal',mu,sigma);  
output; end; keep x; run;
```

Distributions that related to Normal distribution

Certain random variables are related with the Normal distribution and can be generated using those relationships.

- Cauchy distribution: If Y and Z are independent and $Y \sim N(\mu, \sigma^2)$ while $Z \sim N(0, 1)$, then $X = Y/Z$ follows a Cauchy (μ, σ^2) distribution.

- Topic 6
- χ_1^2 distribution: if $Y \sim N(0, 1)$ then $X = Y^2$ follows χ_1^2 distribution.

- χ_n^2 distribution: if Y_1, Y_2, \dots, Y_n are **independent** identically distributed $N(0, 1)$, then

$$X = Y_1^2 + Y_2^2 + \dots + Y_n^2 \sim \chi_n^2$$

- Student's t-distribution with p degrees of freedom t_p : If $Y \sim N(0, 1)$ and $Z \sim \chi_p^2$, then $X = Y/\sqrt{Z/p}$ $\sim t_p$.

10 values $\sim t_{20}$ \rightarrow $rt(10, df=20)$

- F-distribution $F_{m,n}$: if $Y_1 \sim \chi_m^2$, $Y_2 \sim \chi_n^2$ and they are independent, then

$$X = \frac{Y_1/m}{Y_2/n} \sim F_{m,n}$$

Generating χ^2 distribution

We'll generate 10 random numbers that follow χ^2_3 .

- In R.

```
x = rchisq(10, df = 3)
```

- In Python,

```
import numpy as np  
x = np.random.chisquare(df = 3, size = 10)  
print(x)
```

- In SAS.

```
data Chisqgen;  
call streaminit(1234);  
n = 10; df = 3;  
do i = 1 to n;  
x = rand('chisq',df);  
output; end;keep x;run;
```

Generating Binomial Distribution

We'll generate 10 random numbers that follows $Bin(100, p = 0.3)$.

- In R,

```
x = rbinom(n = 10, size = 100 ,prob = 0.3)
```

Note: argument 'size' is the number of trials of the Binomial distribution.

- In Python,

```
import numpy as np
```

```
x = np.random.binomial(n = 100, p = 0.3, size = 10)
```

Note: In Python, argument 'n' is the number of trials of the distribution while argument 'size' is the number of observations that we need to generate.

- In SAS,

```
data Bingen;
```

```
call streaminit(1234);
```

```
p = 0.3; n = 100; *n = number of trials of the distribution;
```

```
do i = 1 to 10; *10 is the number of the observations that we want to  
generate;
```

```
x = rand('binom',p,n);
```

```
output; end;keep x;run;
```

Generating Poisson Distribution

We'll generate 10 random numbers that follow $Poi(3)$, with mean $\lambda = 3$.

- In R,
rpois(10, 3)
- In Python,
import numpy as np
x = np.random.poisson(lam = 3, size = 10)
print(x)
- In SAS,
data Poisgen;
call streaminit(1234);
do i = 1 to 10; *n = 10;
x = rand('poisson',3); *lambda = 3;
output; end;keep x;run;

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- Simulation: Coverage Probability of Confidence Intervals
- Simulation: Properties of Hypothesis Tests
- END

Purposes of using Simulation at our level

- Study properties of estimators. → biasness of an estimator
var of " " "

→ $\bar{X} \rightarrow \mu$
estimator

- Study properties of hypothesis testing procedures. → sig. level α .
type of error → type 1 | error
→ type 2 |
→ power of a test.

What constitutes a simulation study in statistics?


- Simulation involves numerical techniques performed on computers.
- It involves random sampling from probability distributions.



Some issues that matter in Statistics

- Is an estimator biased in finite samples? Is it still consistent under departures from assumptions? What is its sampling variance?

Some issues that matter in Statistics



- Is an estimator biased in finite samples? Is it still consistent under departures from assumptions? What is its sampling variance?
- How does it compare to competing estimators on the basis of bias, precision, etc.?

Some issues that matter in Statistics

- Is an estimator biased in finite samples? Is it still consistent under departures from assumptions? What is its sampling variance?
- How does it compare to competing estimators on the basis of bias, precision, etc.?
- Does a procedure for constructing a confidence interval for a parameter achieve the advertised nominal level of coverage? Like, 95% of the the built 95% CIs cover the true parameter?

Some issues that matter in Statistics

- Is an estimator biased in finite samples? Is it still consistent under departures from assumptions? What is its sampling variance?
- How does it compare to competing estimators on the basis of bias, precision, etc.?
- Does a procedure for constructing a confidence interval for a parameter achieve the advertised nominal level of coverage? Like, 95% of the the built 95% CIs cover the true parameter?
- Does a hypothesis testing procedure attain the advertised level of significance or size?

Some issues that matter in Statistics

- Is an estimator biased in finite samples? Is it still consistent under departures from assumptions? What is its sampling variance?
- How does it compare to competing estimators on the basis of bias, precision, etc.?
- Does a procedure for constructing a confidence interval for a parameter achieve the advertised nominal level of coverage? Like, 95% of the the built 95% CIs cover the true parameter?
- Does a hypothesis testing procedure attain the advertised level of significance or size?
- If it does, what power is possible against different alternatives to the null hypothesis? Do different test procedures deliver different power?

Some issues that matter in Statistics

- Is an estimator biased in finite samples? Is it still consistent under departures from assumptions? What is its sampling variance?
- How does it compare to competing estimators on the basis of bias, precision, etc.?
- Does a procedure for constructing a confidence interval for a parameter achieve the advertised nominal level of coverage? Like, 95% of the the built 95% CIs cover the true parameter?
- Does a hypothesis testing procedure attain the advertised level of significance or size?
- If it does, what power is possible against different alternatives to the null hypothesis? Do different test procedures deliver different power?
- To answer these questions in the absence of analytical results, we can use simulation.


Monte Carlo Simulation Approximation



- When trying to understand a particular parameter, we try to understand the properties of its estimator.

A handwritten diagram below the first bullet point. It shows $\bar{x} \sim$ followed by a box containing $N(\mu; \sigma^2/n)$. To the right of the box is a tilde symbol \sim . An arrow points from the underlined "understand" in the text to the \bar{x} in the diagram.
- The estimator has a true sampling distribution under some conditions about sample (like finite size, etc) and conditions about population (like population follow some specific distribution, etc).
- We want to know this true sampling distribution in order to address the issues mentioned in the previous slide.
- However, very often the true sampling distribution is not possible.
- Hence, we approximate the sampling distribution of an estimator under a particular set of conditions.

Steps of MC Simulation for Approximation

- Generate M independent ^{samples} datasets under the determined conditions. 

Steps of MC Simulation for Approximation

N samples $\rightarrow \bar{x}$ for each sample.
 M

- Generate M independent datasets under the determined conditions.
- Compute the numerical value of the estimator/statistic T from each dataset.
As such, we have T_1, \dots, T_M .

Steps of MC Simulation for Approximation

- Generate M independent datasets under the determined conditions.
- Compute the numerical value of the estimator/statistic T from each dataset. As such, we have T_1, \dots, T_M .
- If M is large enough, the set of simulated estimators T_1, \dots, T_M should be good approximations to the true properties of the estimator or test statistic.

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean ←
- Simulation: Coverage Probability of Confidence Intervals
- Simulation: Properties of Hypothesis Tests
- END

General Ideas

Consider an estimator T for a parameter $\theta = \mu$; $T = \bar{X}$
M samples, each sample \rightarrow get $\bar{X} \Rightarrow M, \bar{X}$

- T_m is the value of T from the m -th dataset, $m = 1, \dots, M$.

General Ideas

Consider an estimator T for a parameter θ .

- T_m is the value of T from the m -th dataset, $m = 1, \dots, M$.
- The mean of all T_m 's is an estimate of the true mean of the sampling distribution of the estimator T .

General Ideas

Consider an estimator T for a parameter θ .

- T_m is the value of T from the m -th dataset, $m = 1, \dots, M$.
- The mean of all T_m 's is an estimate of the true mean of the sampling distribution of the estimator T .
- Simulation allows us to study the properties of the estimator T when the theoretical approach is difficult or untractable.

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1
 - ▶ Sample median, T^2

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1
 - ▶ Sample median, T^2
 - ▶ Sample 10% trimmed mean, T^3

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1
 - ▶ Sample median, T^2
 - ▶ Sample 10% trimmed mean, T^3
- When making comparison, we can base on some criteria like

Examples: Estimators of μ

$$\bar{Y} \rightarrow \mu$$

$$E(\bar{X}) = \mu = 0$$

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1
 - ▶ Sample median, T^2 ←
 - ▶ Sample 10% trimmed mean, T^3 ←
- When making comparison, we can base on some criteria like
 - ▶ Bias(T^k) = $E(T^k) - \mu$, $k = 1, 2, 3$. Estimator with smaller bias is better.

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1
 - ▶ Sample median, T^2
 - ▶ Sample 10% trimmed mean, T^3
- When making comparison, we can base on some criteria like
 - ▶ $\text{Bias}(T^k) = E(T^k) - \mu$, $k = 1, 2, 3$. Estimator with smaller bias is better.
 - ▶ $\text{SD}(T^k) = \sqrt{\text{Var}(T^k)}$, $k = 1, 2, 3$. Estimator with smaller SD is better.

Examples: Estimators of μ

- Say we are interested to compare three estimators for the mean μ of a population distribution based on a random sample Y_1, Y_2, \dots, Y_n . The 3 estimators are:
 - ▶ Sample mean, T^1
 - ▶ Sample median, T^2
 - ▶ Sample 10% trimmed mean, T^3
- When making comparison, we can base on some criteria like
 - ▶ $\text{Bias}(T^k) = E(T^k) - \mu$, $k = 1, 2, 3$. Estimator with smaller bias is better.
 - ▶ $\text{SD}(T^k) = \sqrt{\text{Var}(T^k)}$, $k = 1, 2, 3$. Estimator with smaller SD is better.
 - ▶ $\text{MSE}(T^k) = \text{Bias}(T^k)^2 + (\text{SD}(T^k))^2$, $k = 1, 2, 3$. Smaller MSE is better.

Most common used criteria

Using MSE, compare \bar{X} , sample median, trimmed mean

Estimators of μ : Simulation Procedure

For a particular choice of underlying distribution, μ and sample size n .

- Step 1. Generate independent draws Y_1, Y_2, \dots, Y_n from the distribution.

Estimators of μ : Simulation Procedure

For a particular choice of underlying distribution, μ and sample size n .

- Step 1. Generate independent draws Y_1, Y_2, \dots, Y_n from the distribution.
- Step 2. Compute the value of T^1, T^2, T^3 ^{from} ~~form~~ the sample.

Estimators of μ : Simulation Procedure

For a particular choice of underlying distribution, μ and sample size n .

- Step 1. Generate independent draws Y_1, Y_2, \dots, Y_n from the distribution.
- Step 2. Compute the value of T^1, T^2, T^3 from the sample.
- Repeat {Step 1, Step 2} M times. Hence, we have M copies of each estimator:

sample mean $T^1 : T_1^1, T_2^1, \dots, T_M^1$

sample median $T^2 : T_1^2, T_2^2, \dots, T_M^2$

sample trimmed mean $T^3 : T_1^3, T_2^3, \dots, T_M^3$

M values \Rightarrow mean; var \uparrow MSE

M values \Rightarrow mean; var

M values.
 \Rightarrow mean; var

Estimators of μ : Simulation Procedure

For a particular choice of underlying distribution, μ and sample size n .

- Step 1. Generate independent draws Y_1, Y_2, \dots, Y_n from the distribution.
- Step 2. Compute the value of T^1, T^2, T^3 from the sample.
- Repeat {Step 1, Step 2} M times. Hence, we have M copies of each estimator:

$$T^1 : T_1^1, T_2^1, \dots, T_M^1$$

$$T^2 : T_1^2, T_2^2, \dots, T_M^2$$

$$T^3 : T_1^3, T_2^3, \dots, T_M^3$$

- Calculate the value for criteria (Bias, SD and MSE) for each estimator to make the comparison between the 3 estimators.

Estimators of μ : Comparison

For each $k = 1, 2, 3$, we calculate

- Mean (of M sample means): $\widehat{\text{mean}} = \frac{1}{M} \sum_{m=1}^M T_m^k = \bar{T}^k$
- $\widehat{\text{Bias}}(T^k) = \bar{T}^k - \mu$
- $\widehat{\text{SD}}(T^k) = \sqrt{\frac{1}{M-1} \sum_{m=1}^M (T_m^k - \bar{T}^k)^2}$
- $\widehat{\text{MSE}}(T^k) = \widehat{\text{Bias}}(T^k)^2 + \widehat{\text{SD}}(T^k)^2$.

Depend on what criteria you choose to compare the estimator, you can make conclusion on which estimator is the best to estimate μ .

Estimators of μ : Implementation

- In order to estimate the population mean μ , we consider the 3 estimators: sample mean, sample median and sample 10% trimmed mean.

Estimators of μ : Implementation

- In order to estimate the population mean μ , we consider the 3 estimators: sample mean, sample median and sample 10% trimmed mean.
- Aim: to compare the performance of these 3 estimators.

Estimators of μ : Implementation

- In order to estimate the population mean μ , we consider the 3 estimators: sample mean, sample median and sample 10% trimmed mean.
- Aim: to compare the performance of these 3 estimators.
- Population distribution is $N(170, 10)$ (the true μ is 170).

Estimators of μ : Implementation

- In order to estimate the population mean μ , we consider the 3 estimators: sample mean, sample median and sample 10% trimmed mean.
- Aim: to compare the performance of these 3 estimators.
- Population distribution is $N(170, 10^2)$ (the true μ is 170).
 $sd = 10$.
- Sample size $n = 20$.

Estimators of μ : Implementation

- In order to estimate the population mean μ , we consider the 3 estimators: sample mean, sample median and sample 10% trimmed mean.
- Aim: to compare the performance of these 3 estimators.
- Population distribution is $N(170, 10)$ (the true μ is 170).
- Sample size $n = 20$.
- Simulation size: $M = \underline{10000}$

Estimators of μ : R code (1)

```
> # To compare 3 estimators for location, mu, through a simulation
> # study, the 3 estimators are sample mean, sample median,
> # and 10% trimmed mean.
>
> rm(list= ls())
> M <- 10000 # simulation size (No. of samples)
> n <- 20 # sample size
> mu <- 170 # Mean of the underlying normal distribution
> sd <- 10 # Standard deviation of the underlying normal distribution
> meanx <- numeric(M) # A vector of all sample means
> medx <- numeric(M) # A vector of all sample medians
> trmx <- numeric(M) # A vector of all sample 10% trimmed means
> stdx <- numeric(M) # A vector of all sample standard deviations
> set.seed(1234)
> # Using the same seed number gives same result every time
```

Estimators of μ : R code (2)

```
> for (i in 1:M) {170  
+ x <- rnorm(n,10mu,sd) # Generate a random sample of size n  
+ meanx[i] <- mean(x) # Compute mean of the i-th sample: T^1_i  
+ medx[i] <- median(x) # Compute median of the i-th sample: T^2_i  
+ trmx[i] <- mean(x, trim=0.1)  
+ # Compute the 10% trimmed mean for the i-th sample, i.e. T^3_i  
+ stdx[i] <- sd(x)  
+ # Compute the standard deviation for the i-th sample.  
+ # It is used for computing confidence interval.  
+ }
```

⇒ vector of M sample means! ⚡⚡⚡
_____ median ⚡ u u u
_____ trimmed mean ⚡

Estimators of μ : R code (3)

```
> # Compute mean of M sample means, medians, and trimmed means
> simumean <- apply(cbind(meanx, medx, trmx), 2, mean)
> # "2" in the second argument asks the R to find "means" for
> # all the columns in the matrix given in argument 1.
> # Hence "simumean" is a 1x3 vector consists of
> # the average of the "ns" means, medians, and the 10% trimmed mean
>
> # Compute sd of the M sample means, medians and 10% trimmed means
> simustd <- apply(cbind(meanx, medx, trmx), 2, sd)
> # Compute the bias
> simubias <- simumean - rep(mu, 3)
> # Compute the MSE (Mean Square Error)
> simumse <- simubias2 + simustd2
```

mean(\bar{x}) = 170.
mean(med) = 170.
mean(trim) = 170.

Estimators of μ : R code (4)

```
> estimators <- c("Sample Mean", "Sample Median",  
+                "Sample 10% trimmed mean")  
> # column heading in the output  
> names <- c("True value", "No. of simu", "MC Mean",  
+ "MC Std Deviation", "MC Bias", "MC MSE")  
> # row heading in the output  
> sumdat <- rbind(c(mu, mu, mu), M, simumean, simustd, simubias,  
+                simumse)  
> dimnames(sumdat) <- list(names, estimators)  
> round(sumdat, 4)
```

| | Sample Mean | Sample Median | Sample 10% trimmed mean |
|-------------------------|-------------|---------------|-------------------------|
| <u>True value</u> | 170.0000 | 170.0000 | 170.0000 |
| <u>No. of simu</u> | 10000.0000 | 10000.0000 | 10000.0000 |
| <u>MC Mean</u> | 170.0307 | 170.0314 | 170.0243 |
| <u>MC Std Deviation</u> | 2.2143 | 2.6935 | 2.2757 |
| <u>MC Bias</u> | 0.0307 | 0.0314 | 0.0243 |
| <u>MC MSE</u> | 4.9042 | 7.2559 | 5.1796 |

The sample mean is a better estimator of population mean μ in term of MSE.

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- **Simulation: Coverage Probability of Confidence Intervals**
- Simulation: Properties of Hypothesis Tests
- END

Estimation of μ . Interval Estimate

- The usual 95% confidence interval for μ based on the sample mean from one sample is built as

$$\bar{Y} \pm t_{n-1}(0.025) \times s/\sqrt{n}$$

where s is the standard deviation of that sample.

bc 2 of 95%, $\bar{X} \sim N(\mu; \sigma^2/n)$
 $SD(\bar{X}) = \sigma/\sqrt{n}$
 $SE(\bar{X}) = s/\sqrt{n}$

Estimation of μ : Interval Estimate

- The usual 95% confidence interval for μ based on the sample mean from one sample is built as

$$\bar{Y} \pm t_{n-1}(0.025) \times s/\sqrt{n}$$

where s is the standard deviation of that sample.

- The usual interpretation for 95% CI is: 95% of the confidence intervals constructed as above will cover the true parameter. Is this statement reliable? We'll use simulation to check this statement.

100 → 95 ←
1000 → 950 ←
946
952

Estimation of μ : Interval Estimate

- The usual 95% confidence interval for μ based on the sample mean from one sample is built as

$$\bar{Y} \pm t_{n-1}(0.025) \times s/\sqrt{n}$$

where s is the standard deviation of that sample.

- The usual interpretation for 95% CI is: 95% of the confidence intervals constructed as above will cover the true parameter. Is this statement reliable? We'll use simulation to check this statement.
- For each of the M samples, we'll form a 95% CI for μ . We expect 95% of the M CIs will cover the true mean $\mu = 170$.

Estimation of μ : Interval Estimate

- The usual 95% confidence interval for μ based on the sample mean from one sample is built as

$$\bar{Y} \pm t_{n-1}(0.025) \times s/\sqrt{n}$$

where s is the standard deviation of that sample.

- The usual interpretation for 95% CI is: 95% of the confidence intervals constructed as above will cover the true parameter. Is this statement reliable? We'll use simulation to check this statement.
- For each of the M samples, we'll form a 95% CI for μ . We expect 95% of the M CIs will cover the true mean $\mu = 170$.
- Write R code to get the results of coverage (which should be close to 0.95).

$\bar{X} \pm t \cdot \frac{s}{\sqrt{n}}$; $n=20$.
 CIs of μ : R code
 got M values of sample means ; sd for each sample (\underline{s}) \in
 $\rightarrow M$ CIs for M samples \bar{X} .

```

> # Check the coverage probability of Confidence interval
> # We make use of the data obtained from above simulation study

```

```

> # Get  $t_{\{0.025\}}(0.025)$ 

```

```

> t05 <- qt(0.975, n-1)

```

true value of $\mu = 170$ (slide 42)

```

> d <- 0

```

```

> for (i in 1:M) {

```

```

+ # check if the  $i$ -th CI covers  $\mu$  or not

```

```

+ cover <- (meanx[i] - t05*stdx[i] / sqrt(n) <=  $\mu$ ) &
+           ( $\mu$  <= meanx[i] + t05*stdx[i] / sqrt(n))

```

```

+ d <- d + cover

```

```

+ }

```

```

> coverage <- d/M coverage # this value should close to 0.95

```

```

[1] 0.9538

```

set.seed(1234)

The value of coverage as 0.9538. So the coverage is good.

1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- Simulation: Coverage Probability of Confidence Intervals

- **Simulation: Properties of Hypothesis Tests** ←

- END

Type 1 = $\text{prob}(\text{reject } H_0 \text{ when it is true.})$

$H_0: \mu = \mu_0$ vs

$H_1: \mu \neq \mu_0$

Type 2 = $\text{prob}(\text{do not reject } H_0 \text{ when } H_1 \text{ is true})$

power of a test = $1 - \text{Type 2}$
= $\text{Prob}(\text{reject } H_0 \text{ when } H_1 \text{ is true})$

t-Test for the Mean

$$IQ \sim N(100; \sigma = 10) \leftarrow$$

$$H_0: \mu = 100.$$

$$\bar{X}, S \Rightarrow t \rightarrow \text{small } p\text{-value}$$

Consider the usual t-test for the population mean μ at significance level α :

$$H_0: \mu = \mu_0 \quad \text{vs} \quad H_1: \mu \neq \mu_0$$

Note that, α is the probability of ~~not~~ rejecting H_0 when it is true.

Power of a test is the probability of rejecting H_0 when H_1 is true.

- To evaluate whether the size/level of test can achieve the advertised α :

t-Test for the Mean

Consider the usual t-test for the population mean μ at significance level α :

$$H_0 : \mu = \mu_0 \quad \text{vs} \quad H_1 : \mu \neq \mu_0$$

Note that, α is the probability of not rejecting H_0 when it is true.

Power of a test is the probability of rejecting H_0 when H_1 is true.

- To evaluate whether the **size/level** of test can achieve the advertised α :
 - ▶ Generate data under $H_0 : \mu = \mu_0$ and calculate the proportion of rejections of H_0 .

t-Test for the Mean

Consider the usual t-test for the population mean μ at significance level α :

$$H_0 : \mu = \mu_0 \quad \text{vs} \quad H_1 : \mu \neq \mu_0$$

Note that, α is the probability of not rejecting H_0 when it is true.

Power of a test is the probability of rejecting H_0 when H_1 is true.

- To evaluate whether the **size/level** of test can achieve the advertised α :
 - ▶ Generate data under $H_0 : \mu = \mu_0$ and calculate the proportion of rejections of H_0 .
 - ▶ Get the true probability of rejecting H_0 when it is true (the proportion should be close to α).

t-Test for the Mean

Consider the usual t-test for the population mean μ at significance level α :

$$H_0 : \mu = \mu_0 \quad \text{vs} \quad H_1 : \mu \neq \mu_0$$

Note that, α is the probability of not rejecting H_0 when it is true.

Power of a test is the probability of rejecting H_0 when H_1 is true.

- To evaluate whether the **size/level** of test can achieve the advertised α :
 - ▶ Generate data under $H_0 : \mu = \mu_0$ and calculate the proportion of rejections of H_0 .
 - ▶ Get the true probability of rejecting H_0 when it is true (the proportion should be close to α).
- To estimate the power of the test:

t-Test for the Mean

Consider the usual t-test for the population mean μ at significance level α :

$$H_0 : \mu = \mu_0 \quad \text{vs} \quad H_1 : \mu \neq \mu_0$$

Note that, α is the probability of not rejecting H_0 when it is true.

Power of a test is the probability of rejecting H_0 when H_1 is true.

- To evaluate whether the **size/level** of test can achieve the advertised α :
 - ▶ Generate data under $H_0 : \mu = \mu_0$ and calculate the proportion of rejections of H_0 .
 - ▶ Get the true probability of rejecting H_0 when it is true (the proportion should be close to α).
- To estimate the power of the test:
 - ▶ Generate data under some alternative $H_1 : \mu \neq \mu_0$, say $\mu = \mu_1 \neq \mu_0$ and calculate the proportion of rejections of H_0 .

t-Test for the Mean


Consider the usual t-test for the population mean μ at significance level α :

$$H_0 : \mu = \mu_0 \quad \text{vs} \quad H_1 : \mu \neq \mu_0$$

Note that, α is the probability of not rejecting H_0 when it is true.

Power of a test is the probability of rejecting H_0 when H_1 is true.

- To evaluate whether the **size/level** of test can achieve the advertised α :
 - ▶ Generate data under $H_0 : \mu = \mu_0$ and calculate the proportion of rejections of H_0 .
 - ▶ Get the true probability of rejecting H_0 when it is true (the proportion should be close to α).
- To estimate the power of the test:
 - ▶ Generate data under some alternative $H_1 : \mu \neq \mu_0$, say $\mu = \mu_1 \neq \mu_0$ and calculate the proportion of rejections of H_0 .
 - ▶ Approximate the true probability of rejecting H_0 when H_1 is true.



Checking the Size of t-Test: R code

α

$\rightarrow H_0: \mu = 170$

$\text{data} \sim N(170; \sigma = 10)$

```
> # Check the size of t-test
> # We make use of the data obtained from the above simulation study
> ttests <- (meanx-mu)/(stdx/sqrt(n))
> size <- sum(abs(ttests) > t05)/M
> size
```

[1] 0.0462


test statistic: $t = \frac{\bar{x} - 170}{s/\sqrt{n}}$

1% \rightarrow 5%,
 $t_{0.5} \Rightarrow$ reject H_0
 prob of rejection $\approx 5\%$

We see that the size of 0.05 is approximately being achieved.

$\alpha = 0.05$

$t_{19} = t_{n-1}$



1 Introduction

2 Random Number Generator

3 Simulation Studies in Statistics

- Simulation: Comparing Estimators of Mean
- Simulation: Coverage Probability of Confidence Intervals
- Simulation: Properties of Hypothesis Tests
- END